

Michael C. Hughes

+1-617-371-76 one one • mike@michaelchughes.com
www.michaelchughes.com • michaelchughes

Research Expertise

Machine Learning 2012-

- Unsupervised learning, Semi-supervised learning, Supervised learning
- Probabilistic models: Bayesian nonparametrics, topic models, hidden Markov models, deep generative models
- Posterior estimation methods: variational methods, Markov chain Monte Carlo

Clinical Informatics 2017-

- Tasks: Phenotype discovery, personalized outcome prediction, automated assisted diagnosis
- Data Types: Time series of vitals and lab results from EHR, diagnosis and procedure codes, sociodemographics, imaging (esp. echocardiograms)
- Applications: intensive care early warning systems, suggesting treatments for depression, heart disease diagnosis

Human Activity Analysis 2012-

- Sensor time-series, motion capture, video, images

Education

Brown University
Ph.D., Computer Science. 2016

Brown University
M.S., Computer Science. 2012

Olin College of Engineering
B.S. Electrical & Computer Engineering 2010

Research Experience

Assistant Professor of Computer Science
Tufts University, Medford, MA 2018 - present

- Research on statistical machine learning methods and applications to health informatics.
- Advised Ph.D., M.S., and B.S. students in machine learning research projects.
- Taught advanced undergraduate courses: COMP 135 Intro to ML and COMP 136 Statistical Pattern Recog.
- Developed new course for graduate students with research interests: COMP 150 Bayesian Deep Learning).
- Appointed as the Ann W. Lambertus and Peter Lambertus Assistant Professor in 2019

Postdoctoral fellow: Machine learning to improve clinical decisions in healthcare
Adviser: Prof. Finale Doshi-Velez (Harvard) 2016 - 2018

- Developed semi-supervised models for characterizing and treating depression (with Dr. Perlis and Dr. McCoy).
- Applied time-series models to predict ventilator interventions in the ICU for public dataset of >36,000 patients.
- Created methods for training deep models so they are more interpretable to clinicians or other users.

Postdoc project: Estimating carbon biomass from LiDAR waveforms

Adviser: Prof. Erik Sudderth & Prof. Jim Kellner (Brown U., Ecology & Evolutionary Biology) 2016

- Predicted forest biomass from LiDAR waveforms to better understand land use and climate change.
- Modeled waveforms and biomass predictions jointly via nonparametric regression using our BNPy toolbox.
- Intended for use in NASA's upcoming Global Ecosystem Dynamics Investigation (GEDI).

Ph.D. thesis: Reliable and scalable variational inference for Bayesian nonparametrics

Adviser: Prof. Erik Sudderth 2016

- Thesis Title: Reliable and scalable variational inference for nonparametric mixtures, topics, and sequences
- Developed optimization algorithms for Bayesian nonparametric models that scale to millions of examples.
- Optimized lower bound on marginal likelihood, thus penalizing too simple and too complex explanations.
- Escaped local optima via data-driven proposals that add useful new clusters and remove redundant ones.
- Applied to topic models of 2 million NY Times articles and sequential models of the whole human genome.
- Implemented algorithms in open-source package: Bayesian Nonparametrics for Python (BNPy).

Master's project: Sequential Models for Video and Motion Capture

Adviser: Prof. Erik Sudderth 2012

- Developed methods to discover common actions from many videos of humans performing household exercises.
- Improved existing inference algorithms with data-driven Metropolis-Hastings proposals.

Highlighted Publications

1. "The Tufts fNIRS Mental Workload Dataset & Benchmark for Brain-Computer Interfaces that Generalize." Zhe Huang^d, Liang Wang^d, Giles Blaney, Christopher Slaughter^u, Devon McKeon, Ziyu Zhou, Robert Jacob, and Michael C. Hughes. To appear in Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks, 2021.
2. "Dynamical Wasserstein Barycenters for Time-series Modeling." Kevin C Cheng^d, Shuchin Aeron, Michael C. Hughes, and Eric Miller. To appear in Neural Information Processing Systems (NeurIPS), 2021.
3. "Taming fNIRS-based BCI Input for Better Calibration and Broader Use." Liang Wang^d, Zhe Huang^d, Ziyu Zhou, Devon McKeon, Giles Blaney, Michael C. Hughes, and Robert J. K. Jacob. ACM Symposium on User Interface Software and Technology (UIST), 2021.
4. "A New Semi-supervised Learning Benchmark for Classifying View and Diagnosing Aortic Stenosis from Echocardiograms." Zhe Huang^d, Gary Long, Benjamin Wessler, and Michael C. Hughes. Machine Learning for Healthcare (MLHC), 2021.
5. "Approximate Bayesian Computation for an Explicit-Duration Hidden Markov Model of COVID-19 Hospital Trajectories." Gian Marco Visani^u, Alexandra Hope Lee^b, Cuong Nguyen^m, David M. Kent, John B. Wong, Joshua T. Cohen, and Michael C. Hughes. Machine Learning for Healthcare (MLHC), 2021.
6. "Stochastic Iterative Graph Matching." Linfeng Liu^d, Michael C. Hughes, Soha Hassoun, and Li-Ping Liu. International Conference of Machine Learning (ICML), 2021.
7. "Forecasting COVID-19 Counts At A Single Hospital: A Hierarchical Bayesian Approach." Alexandra Hope Lee^b, Panagiotis Lymperopoulos^m, Joshua T. Cohen, John B. Wong, and Michael C. Hughes. ICLR Workshop on Machine Learning for Preventing and Combating Pandemics, 2021.
8. "Enzyme Promiscuity Prediction Using Hierarchy-Informed Multi-Label Classification." Gian Marco

- Visani^u, Michael C. Hughes, and Soha Hassoun. Bioinformatics, 2021.
9. "On Matched Filtering for Statistical Change Point Detection." Kevin Cheng^d, Eric L Miller, Michael C Hughes, and Shuchin Aeron. IEEE Open Journal of Signal Processing, 2020.
 10. "Assessment of a Prediction Model for Antidepressant Treatment Stability Using Supervised Topic Models." Michael C. Hughes, Melanie F. Pradier, Andrew Slavin Ross, Thomas H. McCoy Jr, Roy H. Perlis, Finale Doshi-Velez. JAMA Network Open, 2020.
 11. "MIMIC-Extract: A Data Extraction, Preprocessing, and Representation Pipeline for MIMIC-III." Shirly Wang^m, Matthew B. A. McDermott, Geeticka Chauhan, Marzyeh Ghassemi, Michael C. Hughes, and Tristan Naumann. CHIL 2020: The ACM Conference on Health, Inference, and Learning, 2020.
 12. "POPCORN: Partially Observed Prediction-Constrained Reinforcement Learning." Joseph Futoma, Michael C. Hughes, and Finale Doshi-Velez. AISTATS, 2020.
 13. "Rapid Model Comparison by Amortizing Across Models." Lily H. Zhang^b, and Michael C. Hughes. Second Symposium on Advances in Approximate Bayesian Inference (AABI 2019), 2019.
 14. "Feature Robustness in Non-stationary Health Records: Caveats to Deployable Model Performance in Common Clinical Machine Learning Tasks." Bret Nestor^d, Matthew B. A. McDermott, Willie Boag, Gabriela Berner, Tristan Naumann, Michael C. Hughes, Anna Goldenberg, and Marzyeh Ghassemi. Machine Learning for Healthcare, 2019.
 15. "Supervised Machine Learning Algorithms Using Patient Related Factors to Predict in-Hospital Mortality Following Acute Myeloid Leukemia Therapy." Nauman Saleem Siddiqui^c, Andreas Klein, Amandeep Godara, Cindy Varga, Rachel J. Buchsbaum, and Michael C. Hughes. Proceedings of 61st Annual Meeting of the American Hematology Society, 2019.
 16. "Semi-Supervised Prediction-Constrained Topic Models." Michael C. Hughes, Gabriel Hope^d, Leah Weiner^d, Thomas H. McCoy Jr, Roy H. Perlis, Erik B. Sudderth, and Finale Doshi-Velez. Artificial Intelligence and Statistics (AISTATS), 2018.
 17. "Beyond Sparsity: Tree Regularization of Deep Models for Interpretability." Mike Wu^u, Michael C. Hughes, Sonali Parbhoo, Maurizio Zazzi, Volker Roth, and Finale Doshi-Velez. Association for Advancement of Artificial Intelligence (AAAI), 2018.
 18. "From Patches to Images: A Nonparametric Generative Model." Geng Ji^d, Michael C. Hughes, and Erik B. Sudderth. International Conference on Machine Learning (ICML), 2017.

Superscripts indicate mentored student's status: u = undergraduate, m = masters, d = Ph.D. student, b = post-bacc, c = medical student. Complete publication list at end of this document.

Highlighted Preprints

1. "Easy Variational Inference for Categorical Observations via a New View of Diagonal Orthant Probit Models." Michael T. Wojnowicz, Shuchin Aeron, Eric Miller, and Michael C. Hughes. Tractable Probabilistic Modeling workshop at UAI, 2021.
2. "Optimizing Clinical Early Warning Systems to Meet False Alarm Constraints." Preetish Rath and Michael C. Hughes. Interpretable Machine Learning for Healthcare (IMLH) workshop at ICML 2021, 2021.
3. "Modeling Graph Node Correlations with Neighbor Mixture Models." Linfeng Liu^d, Michael C. Hughes, and Li-Ping Liu. arXiv, 2021.
4. "Learning Consistent Deep Generative Models from Sparse Data via Prediction Constraints." Gabriel Hope^d, Madina Abdrakhmanova, Xiaoyin Chen, Michael C. Hughes, Erik B. Sudderth. arXiv, 2020.

5. "Prediction-Constrained Training for Semi-Supervised Mixture and Topic Models." Michael C. Hughes, Leah Weiner^d, Gabriel Hope^d, Thomas H. McCoy, Roy H. Perlis, Erik B. Sudderth, and Finale Doshi-Velez. arXiv, 2017.
6. "Fast Learning of Clusters and Topics via Sparse Posteriors." Michael C. Hughes and Erik B. Sudderth. arXiv, 2016.

Honors and Awards

Top 10 Percent Reviewer Award, NeurIPS 2020 2020

- Recognized as one of top 10 percent of more than 3500 expert reviewers at the top international machine learning conference.

Top 400 Reviewer Award, NeurIPS 2019 2019

- Recognized as one of top 400 of more than 3500 expert reviewers at the top international machine learning conference.

Top 200 Reviewer Award, NeurIPS 2018 2018

- Recognized as one of top 200 of more than 3500 expert reviewers at the top international machine learning conference.

Best Paper Award, SoCal NLP Symposium 2018 2018

- Awarded for 2 page summary of our AISTATS 2018 paper.

Nominee for AMIA Clinical Informatics Research Award 2017

- 1 of 7 papers nominated at AMIA's 2017 Joint Summits on Translational Science, out of >50 papers.

NSF Graduate Research Fellowship Award 2011

- Three year award to fund Ph.D. studies. Covers tuition and provides research stipend.

NDSEG Graduate Research Fellowship Award 2011

- Three year funding award. Declined to accept NSF fellowship.

Current Funding Support

Amortized Inference for Large-Scale Graphical Models

NSF CISE: Robust Intelligence: Small 9/1/19 - 8/31/22

- Co-Investigators: Liping Liu (PI, Tufts CS) and Thomas Stopka (Tufts Public Health)
- Total Amount: \$399,923

ACT-NOW: Autonomous Cognitive Technologies for Novelty in Open Worlds

DARPA SAIL-ON Program 11/15/19 - 8/31/22

- SAIL-ON Program: Science of Artificial Intelligence and Learning for Open-world Novelty
- Team at Tufts: Matthias Scheutz (PI, CS), Liping Liu (CS), Jivko Sinapov (CS)
- Team at Arizona State: Chitta Baral (CSE), Subbarao Kambhampati (CSE)

A Benchmark De-identified Echocardiogram Database

Pilot Grant from Tufts CTSI

5/1/21 - 4/30/22

- Full title: A Benchmark De-identified Echocardiogram Database for Studying Automated Diagnoses
- Co-Investigators: Benjamin Wessler (PI)
- Tufts CTSI = Clinical and Translational Science Institute

Estimating the societal value of COVID-19 therapeutics

Tufts Medical Center (Originating Sponsor: Pfizer)

1/1/21 - 12/31/21

- Co-Investigators: Peter Neumann (Tufts CEVR) and Joshua Cohen (Tufts CEVR)

Machine Learning Models for Human Performance Prediction

U.S. Army NSRDEC, Natick, MA (via Tufts CABS)

9/1/20 - 8/31/21

- Full title: Statistical and Machine Learning Models for Data Reduction and Human Performance Prediction
- Co-Investigators: Eric Miller (PI, Tufts ECE) and Shuchin Aeron (Tufts ECE)
- Total Amount: \$496,098
- Part of larger cross-institution MASTR-E project funded by US Army

Past Funding Support

The value of predictive analytics during the COVID epidemic

Tufts Springboard Award (Tufts Univ. Provost's Office)

7/1/20 - 6/30/21

- Full title: Demonstrating the value of a proposed Tufts-led predictive analytics and comparative effectiveness research network during the COVID epidemic
- Co-Investigators: David Kent (Tufts Medical) and Jessica Paulus (Tufts Medical)
- Total Amount: \$50,000

Estimating Individual Treatment Effects

Tufts Collaborates Award (Internal)

7/1/19 - 6/30/20

- Title: 'Estimating Individual Treatment Effects from Randomized Clinical Trials using Machine Learning'
- Co-Investigators: David Kent (Tufts Medical Center)
- Total Amount: \$50,251

Invited Talks

Invited Talk at "I Can't Believe it's not Better" workshop at NeurIPS 2020

12/2020

- Title: I Can't Believe Supervision for Latent Variable Models is not Better: The Case for Prediction constrained training
- Event: I Can't Believe It's Not Better! Workshop (ICBINB at NeurIPS 2020)
- Workshop summary: Bridging the gap between theory and empiricism in probabilistic machine learning
- Talk summary: Makes case for our recent work on prediction constrained training, from AISTATS 2018, AISTATS 2020, and in preparation work

Invited Talk at Northwell Health ML group 07/2020

- Title: Optimizing Machine Learning Models for Interpretable, Actionable Predictions on Electronic Health Records
- Event: Regular virtual meeting of a ML research working group at Northwell Health (large healthcare provider in NYC)
- Summarizes recent MLHC 2019, CHIL 2020, AISTATS 2020, and JAMA Netw. Open 2020 papers

Invited Talk at U. Arizona 02/2020

- Title: Overcoming model misspecification in structured clustering and reinforcement learning with prediction constrained training
- Event: Regular meeting of a research working group at U. Arizona funded by NSF TRIPODS award
- Summarizes our recent AISTATS 2020 paper

Invited Short Talk at Duke Clinical Research Institute Think Tank meeting 01/2020

- Talk: Preferred Quality Metrics for Clinical Prediction Models
- Event: Leveraging Artificial Intelligence and Machine Learning Methods and Approaches to Transform Clinical Trial Design, Planning, and Execution
- Host Organization: Duke Clinical Research Institute
- An invitation-only event in Washington D.C. gathering 40 ML experts, clinical experts, and policy makers

Invited Talk at Meeting of Critical Care Directors in Madrid, Spain 01/2020

- Talk Title: Optimizing Machine Learning Models for Interpretable, Actionable Predictions
- Event: Reunión Sobre Nuevas Tecnologías en el Tratamiento de Datos Clínicos Electrónicos
- Translation: Meeting on New Technologies for Processing Electronic Health Records
- Hosts: RGI Informatics (Dr. Richard Goldstein, CEO) and Fuenlabrada University Hospital (Joaquín Álvarez, head of ICU).
- An invitation-only hosted event in Madrid for 30 directors of intensive care units around Madrid, Spain

Invited Mentor at 2019 PLA General Hospital - MIT Critical Data Datathon 11/2019

- 4th annual PLAGH-MIT Datathon
- Event held in Beijing, China with 25 teams of local clinicians and computational scientists
- Team goal: Answer clinical question on historical public dataset (MIMIC) over 1 weekend
- Event goal: Develop local teams' skills via intense practice with expert oversight
- My role: Advise teams toward principled and clinically-useful analysis

Invited Talk at BNP 2019 06/2019

- Title: Scalable and Reliable Variational Inference for Dirichlet Process Clustering with Sparse Assignments
- Venue: 12th International Conference on Bayesian Nonparametrics
- Summarizes the effective learning methods behind our BNPy toolbox

Invited Tutorial at MLHC 2018 08/2018

- Machine Learning for Clinicians: Advances for Multi-Modal Health Data at MLHC '18
- Designed to help clinicians understand enough modern machine learning to collaborate successfully with ML researchers.

Invited Panelist 12/2016

- Software panel at Advances in Approximate Bayesian Inference workshop at NIPS '16.

Professional Service

Area Chair

- 2021 - ACM CHIL and MLHC

Senior Program Committee / Meta-Reviewer

- 2021 - AAAI
- 2020 - AAAI

Program Committee / Reviewer

- 2020 - NeurIPS, AISTATS, ICLR, MLHC
- 2019 - NeurIPS (reviewer award), AISTATS, ICLR
- 2018 - NeurIPS (reviewer award), AAAI, AISTATS, ICLR, AMIA CRI
- 2017 - NeurIPS, ICML, AAAI
- 2016 - NeurIPS
- 2015 - NeurIPS, ICML
- 2014 - NeurIPS, ICML
- 2013 - NeurIPS (reviewer award)

Workshop Organizer: ML4H at NeurIPS 2018 2018

- Machine Learning for Health workshop at NeurIPS '18 (NeurIPS ML4H 2018).
- Full-day workshop with invited keynotes, accepted papers/posters, and lively panel discussions.
- Provided a forum for interdisciplinary interaction between clinicians, statisticians, and computer scientists.
- Helped with website, PR, and continuity in peer-review process from previous years.

Workshop Organizer: BNP at NeurIPS 2018 2018

- All of Bayesian Nonparametrics workshop at NeurIPS '18 (NeurIPS BNP 2018).
- Full-day workshop with invited keynotes, accepted papers/posters, and lively panel discussions.
- Helped with peer-review process for accepted posters, <https://sites.google.com/view/nipsbnp2018/schedule>.

Workshop Organizer: ML4H at NeurIPS 2017 2017

- Machine Learning for Health workshop at NIPS '17 (NIPS ML4H 2017)
- Full-day workshop with invited keynotes and panels involving clinicians, statisticians, and computer scientists.
- Organized peer-review process for 118 submitted papers.

Workshop Organizer: BNP at NeurIPS 2016 2016

- Practical Bayesian Nonparametrics workshop at NIPS '16.
- Full-day workshop with invited speakers, contributed talks, two panel discussions, and lively poster session.
- Led decisions on >25 submitted papers based on peer review.

Teaching and Mentorship

Tufts CS Dept.

Course: COMP 136 Statistical Pattern Recognition Spring 2021

- Taught advanced statistical learning course to 29 students
- Course format: 5 math-intensive homeworks, 5 coding-intensive homeworks, 5 short quizzes, 2 exams

Tufts CS Dept.

Course: COMP 135 Intro to Machine Learning

Fall 2020

- Taught core principles of machine learning to 95 students
- Course format: 3 open-ended projects, 5 homeworks (conceptual and code questions), and 5 quizzes

Tufts CS Dept.

Course: COMP 136 Statistical Pattern Recognition

Spring 2020

- Taught advanced statistical learning course to 35 students
- Course format: 5 math-intensive homeworks, 5 coding-intensive homeworks, 5 short quizzes, 2 exams

Tufts CS Dept.

Course: COMP 150 Bayesian Deep Learning

Fall 2019

- Taught advanced topics seminar to 23 students
- Course format: weekly homeworks for first month, then 2-month open-ended team project

Tufts CS Dept.

Course: COMP 135 Introduction to Machine Learning

Spring 2019

- Taught core principles of machine learning to about 50 students
- Course format: 3 open-ended projects, weekly homeworks, and 2 exams

Tufts CS Dept.

Course: COMP 150 Bayesian Deep Learning

Fall 2018

- Taught advanced topics seminar to about 18 students
- Course format: weekly homeworks for first month, then 2-month open-ended team project
- One project resulted in publication at IEEE conference (ICDL-EpiRob 2019)

Tufts University

Research Mentor

2018-

- Mentored masters student projects:
 - – Yu Liu. 2020-21 CS DS thesis: An Evaluation Pipeline for Heterogeneous Treatment Effect Prediction
 - – Xi Chen. 2020 MS project: Bayesian Nonparametric Mixture Models for Missing Data
- Mentored projects for post-bacc students:
 - – Ally Lee. 2020: Bayesian Analysis of Autoregressive Models for Multi-Site Hospital Admission Forecasting
 - – Lily H. Zhang 2019-20: Any Parameter Encoders for Topic Models: Variational Encoders that amortize across models as well as data
- Mentored undergraduate projects:
 - – Hezekiah Branch. 2020 Tufts LSAMP fellow: Supervised Learning for Clinical Multivariate Time-series
 - – Manh Duc Nguyen. 2019 CS honors thesis: Particle-based algorithms for Bayesian Neural Networks - Hamiltonian Monte Carlo and Stein Variational Gradient Descent

Harvard University SEAS

Research Mentor

2016-2017

- Mentored undergraduate senior thesis projects on Bayesian nonparametric inference.
- Frederick Widjaja. 2017 honors thesis: Streaming Variational Inference for the Indian Buffet Process.
- Madhu Vijay. 2017 honors thesis: Characterizing Posterior Uncertainty for the Indian Buffet Process.

Brown University

Research Mentor

2014-2016

- o Mentored students on projects related to Bayesian nonparametric clustering and the BNPy Python package.
- o William Stephenson. 2015 undergraduate honors thesis: Variational Inference for Hierarchical Dirichlet Process based Nonparametric Models.
- o Sonia Phene. 2015 undergraduate honors thesis: Multiprocessor Parallelization of Variational Inference for Bayesian Nonparametric Topic Models.
- o Mengrui Ni. 2015 masters project: Variational Inference for Beta-Bernoulli Dirichlet Process Mixture Models.
- o Mert Terzihan. 2015 masters project.

Lead Graduate TA for CS 142: Intro to Machine Learning

Fall 2013

- o Led weekly 1 hour recitation session to review key concepts for 50+ students.
- o Designed homework assignments and exam questions.

Outreach Experience

Tufts DIAMONDS Program

Medford, MA

Research Mentor

2021

- o Mentored two students in data science research projects.

TEALS and Boston Latin Academy

Roxbury, MA

Volunteer AP Computer Science Instructor

2014-2016

- o Taught 1-2 classes / week for 2 years via TEALS "CS in every high school" initiative sponsored by Microsoft.
- o Developed hands-on lessons to excite students from diverse backgrounds about computational thinking.
- o Mentored full-time teacher Ingrid Roche as she transitioned from media arts to AP computer science (Java).

Harvard Humanitarian Initiative

Cambridge, MA

Signal Program Fellow

2014

- o Developed prototype detector for common housing structures in sub-Saharan Africa from satellite images.
- o Intended for humanitarian oversight of conflict areas where burning structures is common attack pattern.
- o Featured in TEDx talk: <https://youtu.be/u719rBw0nwU>

Olin College Engineering Discovery

Needham, MA

Co-Founder and Curriculum Director

2007-2010

- o Managed 15 undergrads in developing hands-on lessons for 4th-8th graders.
- o Hosted workshops for 30 children to design, build, and launch bottle rockets.
- o Pioneered green energy workshop which earned over \$750 in outside funding.

Industry Experience

Google

Mountain View, CA

Software Engineering Intern

Summer 2013

- o Improved walking/biking/running classifier using smartphone accelerometer data.
- o Led collection of dataset from dozens of individuals for classifier evaluation via custom Android app.

All Conference Publications (in reverse chronological order)

1. "The Tufts fNIRS Mental Workload Dataset & Benchmark for Brain-Computer Interfaces that Generalize." Zhe Huang^d, Liang Wang^d, Giles Blaney, Christopher Slaughter^u, Devon McKeon, Ziyu Zhou, Robert Jacob, and Michael C. Hughes. To appear in Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks, 2021.
2. "Dynamical Wasserstein Barycenters for Time-series Modeling." Kevin C Cheng^d, Shuchin Aeron, Michael C. Hughes, and Eric Miller. To appear in Neural Information Processing Systems (NeurIPS), 2021.
3. "Taming fNIRS-based BCI Input for Better Calibration and Broader Use." Liang Wang^d, Zhe Huang^d, Ziyu Zhou, Devon McKeon, Giles Blaney, Michael C. Hughes, and Robert J. K. Jacob. ACM Symposium on User Interface Software and Technology (UIST), 2021.
4. "A New Semi-supervised Learning Benchmark for Classifying View and Diagnosing Aortic Stenosis from Echocardiograms." Zhe Huang^d, Gary Long, Benjamin Wessler, and Michael C. Hughes. Machine Learning for Healthcare (MLHC), 2021.
5. "Approximate Bayesian Computation for an Explicit-Duration Hidden Markov Model of COVID-19 Hospital Trajectories." Gian Marco Visani^u, Alexandra Hope Lee^b, Cuong Nguyen^m, David M. Kent, John B. Wong, Joshua T. Cohen, and Michael C. Hughes. Machine Learning for Healthcare (MLHC), 2021.
6. "Stochastic Iterative Graph Matching." Linfeng Liu^d, Michael C. Hughes, Soha Hassoun, and Li-Ping Liu. International Conference of Machine Learning (ICML), 2021.
7. "MIMIC-Extract: A Data Extraction, Preprocessing, and Representation Pipeline for MIMIC-III." Shirley Wang^m, Matthew B. A. McDermott, Geeticka Chauhan, Marzyeh Ghassemi, Michael C. Hughes, and Tristan Naumann. CHIL 2020: The ACM Conference on Health, Inference, and Learning, 2020.
8. "POPCORN: Partially Observed Prediction-Constrained Reinforcement Learning." Joseph Futoma, Michael C. Hughes, and Finale Doshi-Velez. AISTATS, 2020.
9. "Optimal Transport Based Change Point Detection and Time Series Clustering." Kevin Cheng^d, Shuchin Aeron, Michael C. Hughes, Erika Hussey, and Eric Miller. IEEE ICASSP 2020, 2020.
10. "Regional Tree Regularization for Interpretability in Deep Neural Networks." Mike Wu^d, Sonali Parbhoo, Michael C. Hughes, Ryan Kindle, Leo Celi, Maurizio Zazzi, Volker Roth, and Finale Doshi-Velez. AAAI, 2020.
11. "Feature Robustness in Non-stationary Health Records: Caveats to Deployable Model Performance in Common Clinical Machine Learning Tasks." Bret Nestor^d, Matthew B. A. McDermott, Willie Boag, Gabriela Berner, Tristan Naumann, Michael C. Hughes, Anna Goldenberg, and Marzyeh Ghassemi. Machine Learning for Healthcare, 2019.
12. "Supervised Machine Learning Algorithms Using Patient Related Factors to Predict in-Hospital Mortality Following Acute Myeloid Leukemia Therapy." Nauman Saleem Siddiqui^c, Andreas Klein, Amandeep Godara, Cindy Varga, Rachel J. Buchsbaum, and Michael C. Hughes. Proceedings of 61st Annual Meeting of the American Hematology Society, 2019.
13. "Sensorimotor Cross-Behavior Knowledge Transfer for Grounded Category Recognition." Gyan Tatiya^d, Ramtin Hosseini^d, Michael C. Hughes, and Jivko Sinapov. Joint IEEE International Conference on Development and Learning and Epigenetic Robotics (ICDL-EpiRob), 2019.
14. "Semi-Supervised Prediction-Constrained Topic Models." Michael C. Hughes, Gabriel Hope^d, Leah Weiner^d, Thomas H. McCoy Jr, Roy H. Perlis, Erik B. Sudderth, and Finale Doshi-Velez. Artificial

- Intelligence and Statistics (AISTATS), 2018.
15. "Beyond Sparsity: Tree Regularization of Deep Models for Interpretability." Mike Wu^u, Michael C. Hughes, Sonali Parbhoo, Maurizio Zazzi, Volker Roth, and Finale Doshi-Velez. Association for Advancement of Artificial Intelligence (AAAI), 2018.
 16. "From Patches to Images: A Nonparametric Generative Model." Geng Ji^d, Michael C. Hughes, and Erik B. Sudderth. International Conference on Machine Learning (ICML), 2017.
 17. "Right for the Right Reasons: Training Differentiable Models by Constraining their Explanations." Andrew Slavin Ross^m, Michael C. Hughes, and Finale Doshi-Velez. International Joint Conference on Artificial Intelligence (IJCAI), 2017.
 18. "Predicting Intervention Onset in the ICU with Switching State Space Models." Marzyeh Ghassemi, Mike Wu^u, Michael C. Hughes, Peter Szolovits, and Finale Doshi-Velez. AMIA Summit on Clinical Research Informatics, 2017.
 19. "Scalable Adaptation of State Complexity for Nonparametric Hidden Markov Models." Michael C. Hughes, William Stephenson^u, and Erik B. Sudderth. Neural Information Processing Systems (NIPS), 2015.
 20. "Reliable and Scalable Variational Inference for the Hierarchical Dirichlet Process." Michael C. Hughes, Dae Il Kim, and Erik B. Sudderth. Artificial Intelligence & Statistics (AISTATS), 2015.
 21. "Memoized Online Variational Inference for Dirichlet Process Mixture Models." Michael C. Hughes and Erik B. Sudderth. Neural Information Processing Systems (NIPS), 2013.
 22. "Effective Split-Merge Monte Carlo Methods for Nonparametric Models of Sequential Data." Michael C. Hughes, Emily Fox, and Erik B. Sudderth. Neural Information Processing Systems (NIPS), 2012.
 23. "The Nonparametric Metadata Dependent Relational Model." Dae Il Kim, Michael C. Hughes, and Erik B. Sudderth. International Conference on Machine Learning (ICML), 2012.

All Journal Publications (in reverse chronological order)

1. "The role of machine learning in clinical research: transforming the future of evidence generation." E. Hope Weissler, Tristan Naumann, Tomas Andersson, Rajesh Ranganath, Olivier Elemento, Yuan Luo, Daniel F. Freitag, James Benoit, Michael C. Hughes, Faisal Khan, Paul Slater, Khader Shameer, Matthew Roe, Emmette Hutchison, Scott H. Kollins, Uli Broedl, Zhaoling Meng, Jennifer L. Wong, Lesley Curtis, Erich Huang and Marzyeh Ghassemi. *Trials*, 2021.
2. "Optimizing for Interpretability in Deep Neural Networks with Simulable Decision Trees." Mike Wu^d, Sonali Parbhoo, Michael C. Hughes, Volker Roth, and Finale Doshi-Velez. *Journal of Artificial Intelligence Research (JAIR)*, 2021.
3. "Enzyme Promiscuity Prediction Using Hierarchy-Informed Multi-Label Classification." Gian Marco Visani^u, Michael C. Hughes, and Soha Hassoun. *Bioinformatics*, 2021.
4. "On Matched Filtering for Statistical Change Point Detection." Kevin Cheng^d, Eric L. Miller, Michael C. Hughes, and Shuchin Aeron. *IEEE Open Journal of Signal Processing*, 2020.
5. "A Framework for Sensorimotor Cross-Perception and Cross-Behavior Knowledge Transfer for Object Categorization." Gyan Tatiya^d, Ramtin Hosseini^d, Michael C. Hughes, and Jivko Sinapov. *Frontiers in Robotics and AI*, 2020.
6. "Assessment of a Prediction Model for Antidepressant Treatment Stability Using Supervised Topic Models." Michael C. Hughes, Melanie F. Pradier, Andrew Slavin Ross, Thomas H. McCoy Jr, Roy H. Perlis, Finale Doshi-Velez. *JAMA Network Open*, 2020.

7. "Predicting change in diagnosis from major depression to bipolar disorder after antidepressant initiation." Melanie F. Pradier, Michael C. Hughes, Thomas H. McCoy Jr, Sergio A. Barroilhet, Finale Doshi-Velez, and Roy H. Perlis. *Neuropsychopharmacology*, 2020.
8. "Predicting Treatment Discontinuation after Antidepressant Initiation." Melanie F. Pradier, Thomas H. McCoy, Michael C. Hughes, Roy H. Perlis, and Finale Doshi-Velez. *Translational Psychiatry*, 2020.
9. "Refinery: An Open Source Topic Modeling Web Platform." Daeil Kim, Benjamin F. Swanson, Michael C. Hughes, and Erik B. Sudderth. *JMLR Machine Learning Open Source Software (MLOSS)*, 2017.
10. "Joint Modeling of Multiple Time Series via the Beta Process with Application to Motion Capture Segmentation." Emily Fox, Michael C. Hughes, Erik B. Sudderth, and Michael I. Jordan. *Annals of Applied Statistics*, Vol. 8(3), 2014.

All Workshop Papers (in reverse chronological order)

1. "Easy Variational Inference for Categorical Observations via a New View of Diagonal Orthant Probit Models." Michael T. Wojnowicz, Shuchin Aeron, Eric Miller, and Michael C. Hughes. *Tractable Probabilistic Modeling workshop at UAI*, 2021.
2. "Optimizing Clinical Early Warning Systems to Meet False Alarm Constraints." Preetish Rath and Michael C. Hughes. *Interpretable Machine Learning for Healthcare (IMLH) workshop at ICML 2021*, 2021.
3. "Prediction-Constrained Hidden Markov Models for Semi-Supervised Classification." Gabriel Hope^d, Michael C. Hughes, Finale Doshi-Velez, and Erik B. Sudderth. *Time Series Workshop at ICML*, 2021.
4. "Evaluating the Use of Reconstruction Error for Novelty Localization." Patrick Feeney^d and Michael C. Hughes. *Uncertainty and Robustness in Deep Learning (UDL) workshop at ICML 2021*, 2021.
5. "Forecasting COVID-19 Counts At A Single Hospital: A Hierarchical Bayesian Approach." Alexandra Hope Lee^b, Panagiotis Lymperopoulos^m, Joshua T. Cohen, John B. Wong, and Michael C. Hughes. *ICLR Workshop on Machine Learning for Preventing and Combating Pandemics*, 2021.
6. "Using Hierarchy-Informed Multi-Label Classification for Enzyme Promiscuity Prediction." Gian Marco Visani^u, Michael C. Hughes, and Soha Hassoun. *Machine Learning in Computational Biology Workshop (MLCB)*, 2020.
7. "Rapid Model Comparison by Amortizing Across Models." Lily H. Zhang^b, and Michael C. Hughes. *Second Symposium on Advances in Approximate Bayesian Inference (AABI 2019)*, 2019.
8. "Classification of Enzyme Promiscuity Using Positive, Unlabeled, and Hard Negative Examples." Gian Marco Visani, Michael C. Hughes and Soha Hassoun. *Machine Learning in Computational Biology Workshop (MLCB)*, 2019.
9. "Prediction-Constrained POMDPs." Joseph Futoma, Michael C. Hughes, and Finale Doshi-Velez. *Reinforcement Learning under Partial Observability (RLPO) workshop at NeurIPS 2018*, 2018.
10. "Rethinking clinical prediction: Why machine learning must consider year of care and feature aggregation." Bret Nestor^d, Matthew B. A. McDermott, Geeticka Chauhan, Tristan Naumann, Michael C. Hughes, Anna Goldenberg, Marzyeh Ghassemi. *Machine Learning for Healthcare (ML4H) workshop at NeurIPS 2018*, 2018.
11. "Prediction-Constrained Topic Models for Antidepressant Prediction." Michael C. Hughes, Gabriel Hope^d, Leah Weiner^d, Thomas H. McCoy, Roy H. Perlis, Erik B. Sudderth, and Finale Doshi-Velez. *NIPS Workshop on Machine Learning for Health (NIPS ML4H)*, 2017.

12. "Associations between aboveground forest biomass and waveform LiDAR metrics: implications for modeling footprint-level biomass using Global Ecosystem Dynamics Investigation data." J. Kellner, J. B. Blair, L. Duncanson, L., S. Hancock, M. A. Hofton, M. C. Hughes, S. Marselis, S., J. Armston, E. B. Sudderth, H. Tang, L. Weiner^d, and R. Dubayah. American Geophysical Union, Fall General Assembly, 2016.
13. "Supervised topic models for clinical interpretability." Michael C. Hughes, Huseyin Melih Elibol, Thomas McCoy, Roy Perlis, and Finale Doshi-Velez. NIPS Workshop on Machine Learning for Health (NIPS ML4H), 2016.
14. "BNPy: Reliable and scalable variational inference for Bayesian nonparametric models." Michael C. Hughes and Erik B. Sudderth. 3rd NIPS Workshop on Probabilistic Programming, 2014.
15. "Nonparametric Discovery of Activity Patterns from Video Collections." Michael C. Hughes and Erik B. Sudderth. CVPR Workshop on Perceptual Organization in Computer Vision (POCV), 2012.